


Introduction to Ray for Distributed Applications

Jules Damji, Anyscale
@2twitme



\$whoami (Jules)

- Lead Developer Advocate @Anyscale
- Senior Developer Advocate @Databricks
- Led Developer Advocacy @Hortonwork
- Held SWE positions:
 - + Sun Microsystems
 - + Netscape
 - + @Home
 - + Loudcloud/Opsware
 - + Verisign

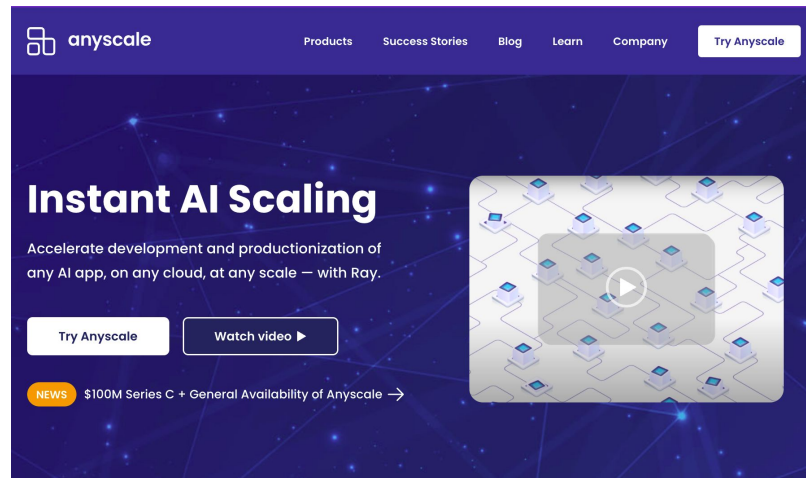


Anyscale

Who we are: Original creators of Ray

What we do: Managed compute platform to scale AI & Python workloads

Why do it: Scaling is a necessity, scaling is hard; make distributed computing easy and simple for everyone



Agenda

- Why & What's Ray & Ray Ecosystem
- Ray Architecture & Components
- Ray Core Design Patterns & APIs
- Modules [1 - 3]
- Closing Q & A
- Happy Hour 🍺 + Meetup

Why Ray


- Machine learning is pervasive
- Distributed computing is a necessity
- Python is the default language for DS/ML

Blessings of scale ...

The blessings of scale


AI training runs, estimated computing resources used

Floating-point operations, selected systems, by type, log scale




Sources: "Compute trends across three eras of machine learning", by J. Sevilla et al., arXiv, 2022; Our World in Data


Compute - supply demand problem



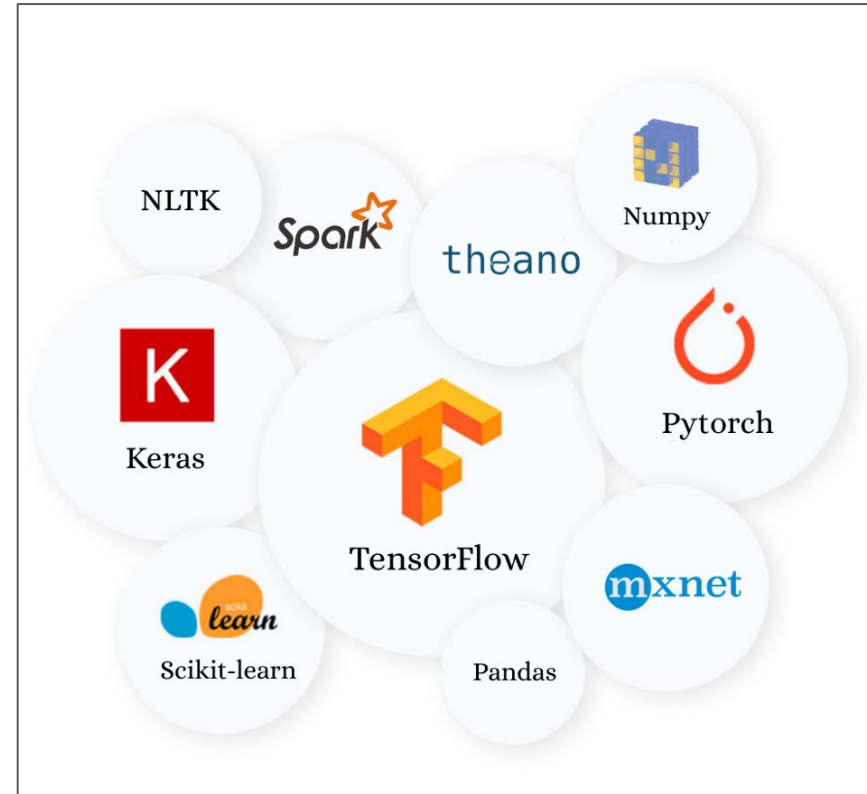
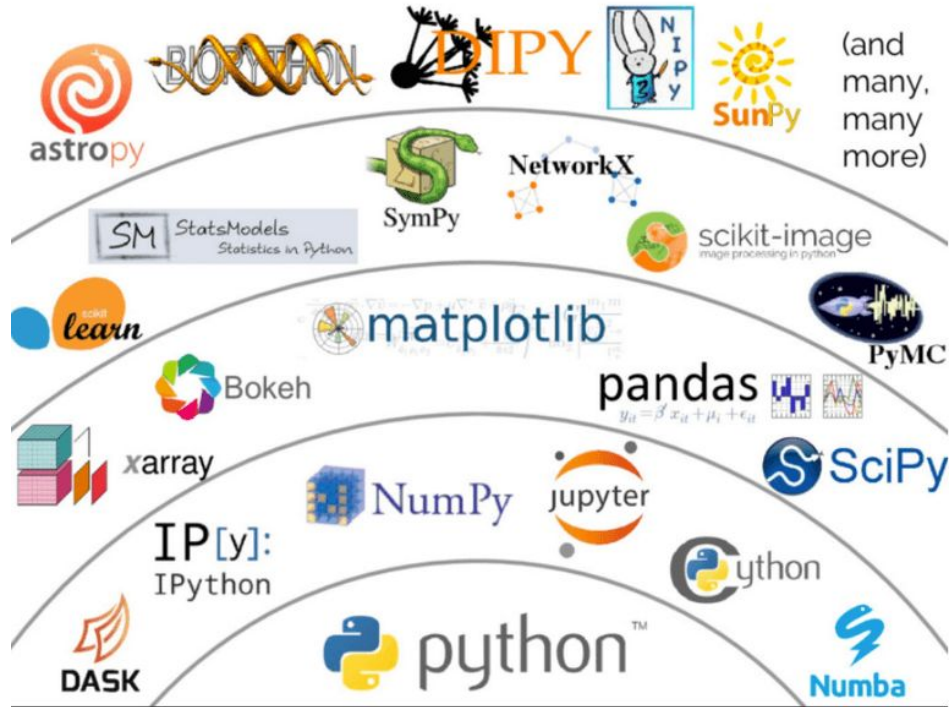
Specialized hardware is not enough



No way out but to distribute!



Python data science ecosystem

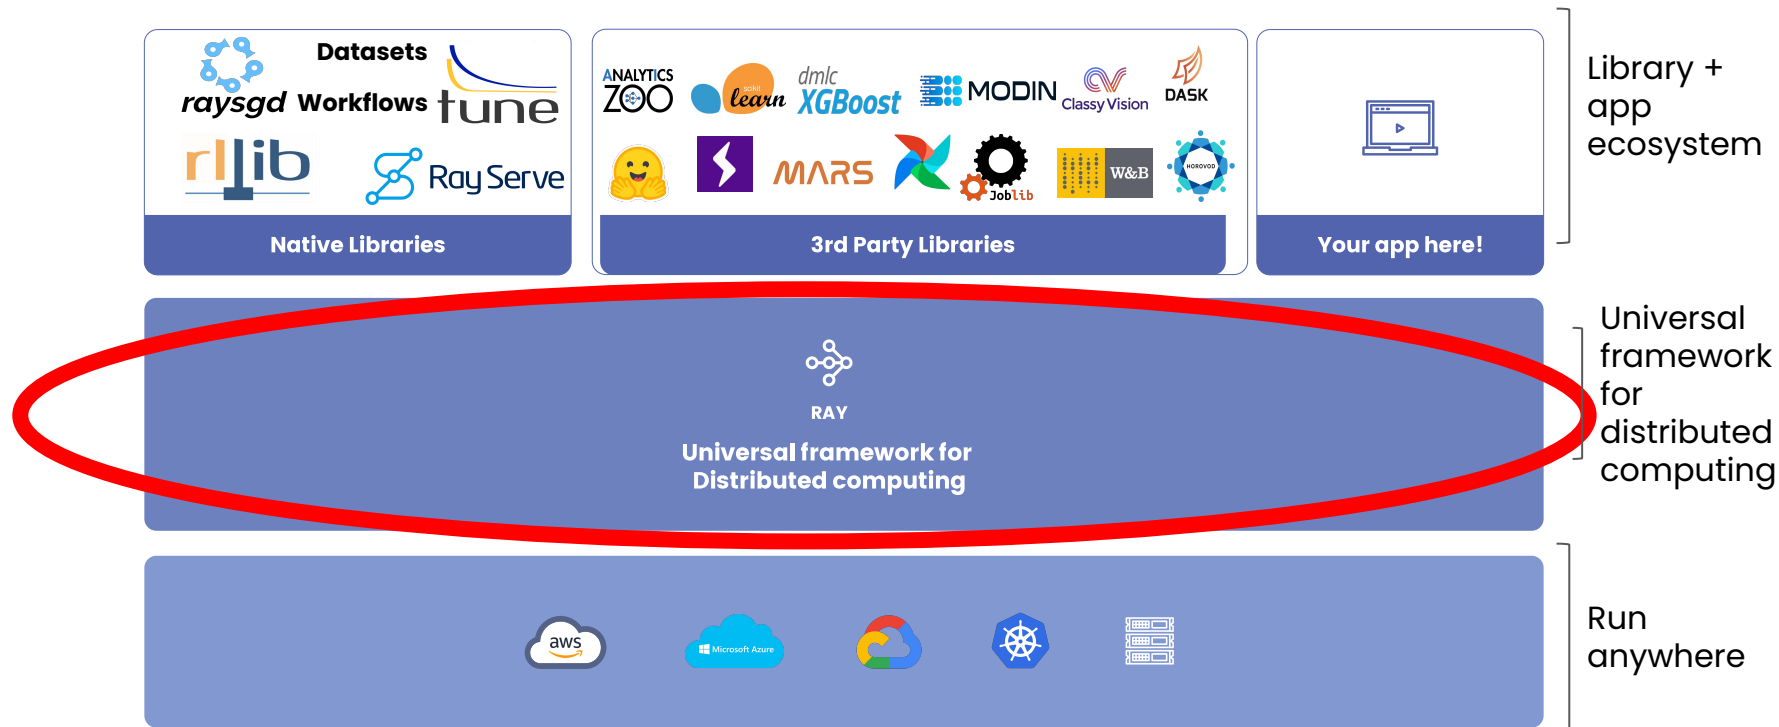


What is Ray

- A simple/general-purpose library for distributed computing
- An ecosystem of Python libraries (for scaling ML and more)
- Runs on laptop, public cloud, K83, on-premise

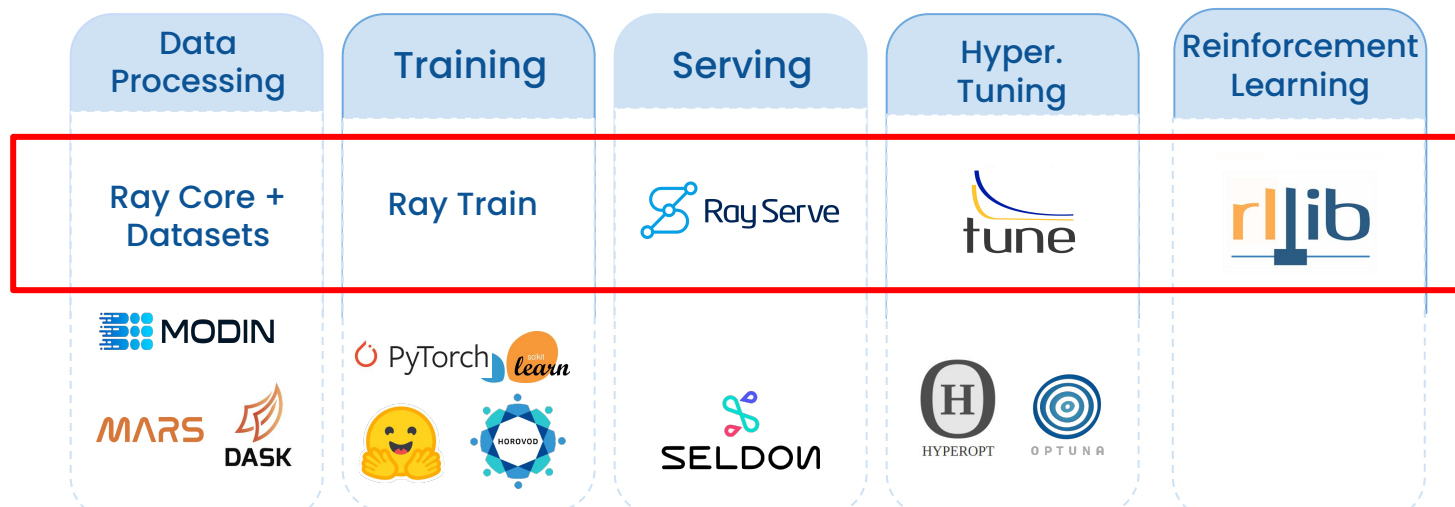
A layered cake of functionality and capabilities for scaling ML workloads

A Layered Cake and Ecosystem



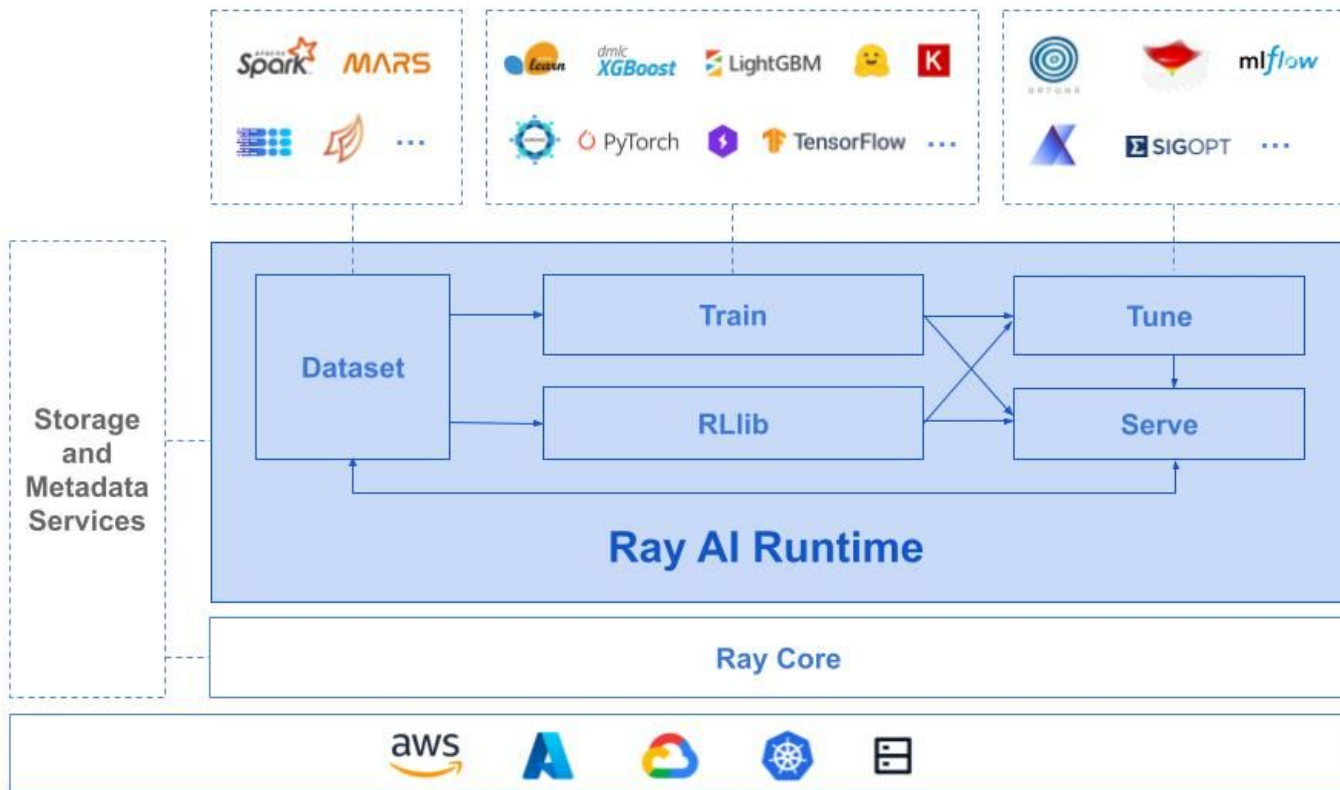
Rich ecosystem of scaling ML workloads

Built-in
“batteries
included”
libraries



Only use the libraries you need!

Ray AI Runtime (Ray AIR)



Ray 2.0 & Ray AIR sessions

- Introduction Ray AI Runtime
- State of Ray Serve in 2.0
- Shuffling 100TB with Ray Datasets
- Ray Observability: Present & future
- Many others in Ray Deep Dives track ...

Who's using it

amazon



Uber




McKinsey
& Company



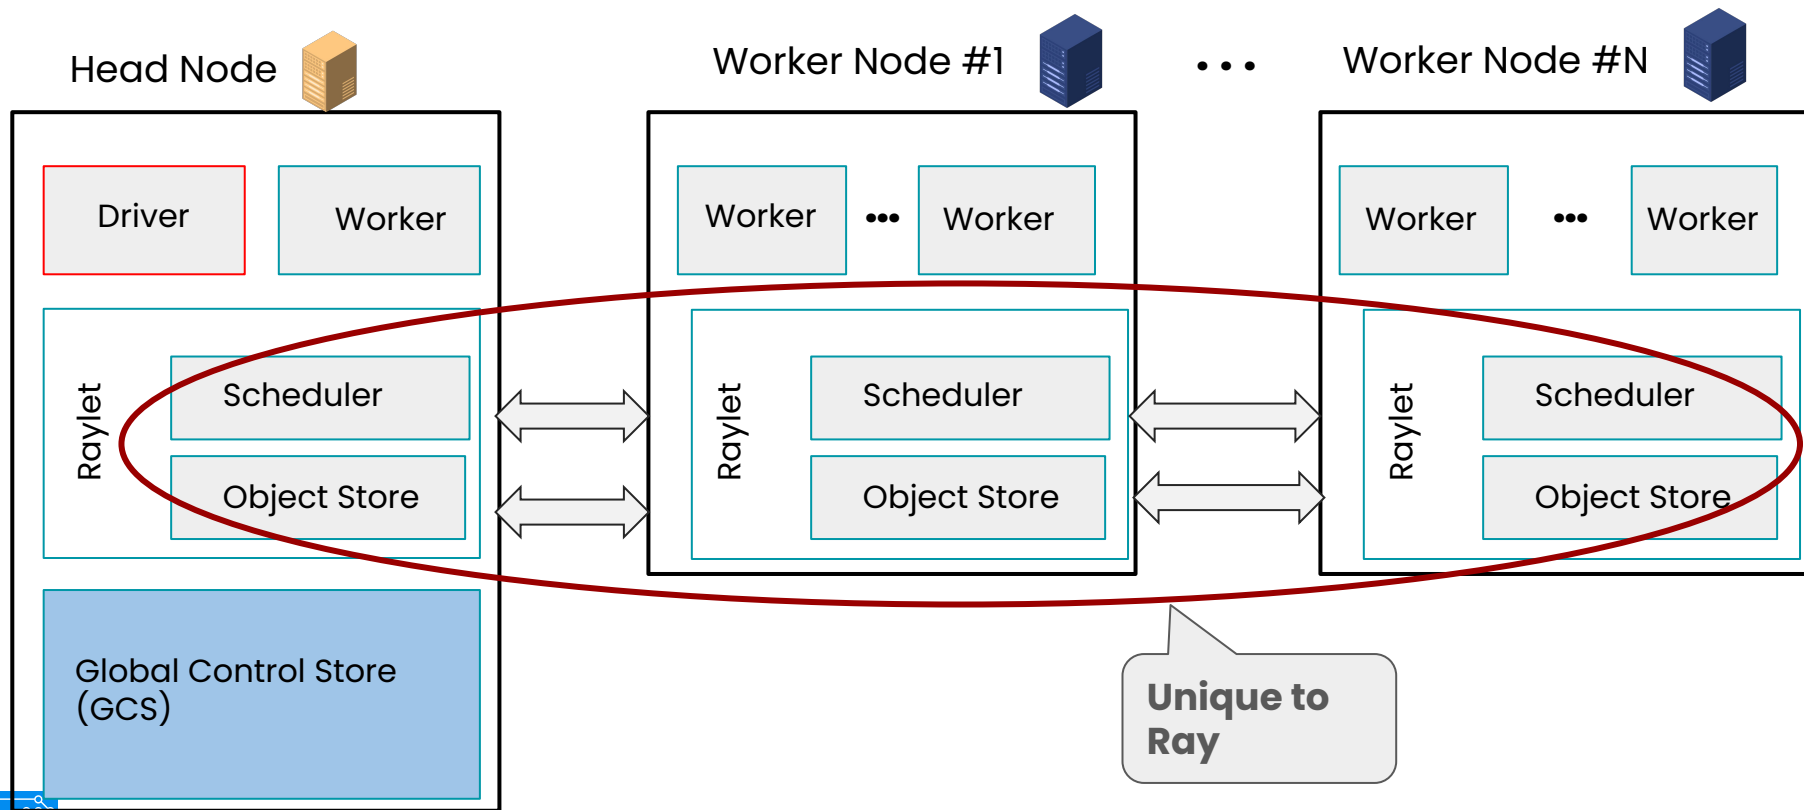
**WILD
LIFE**

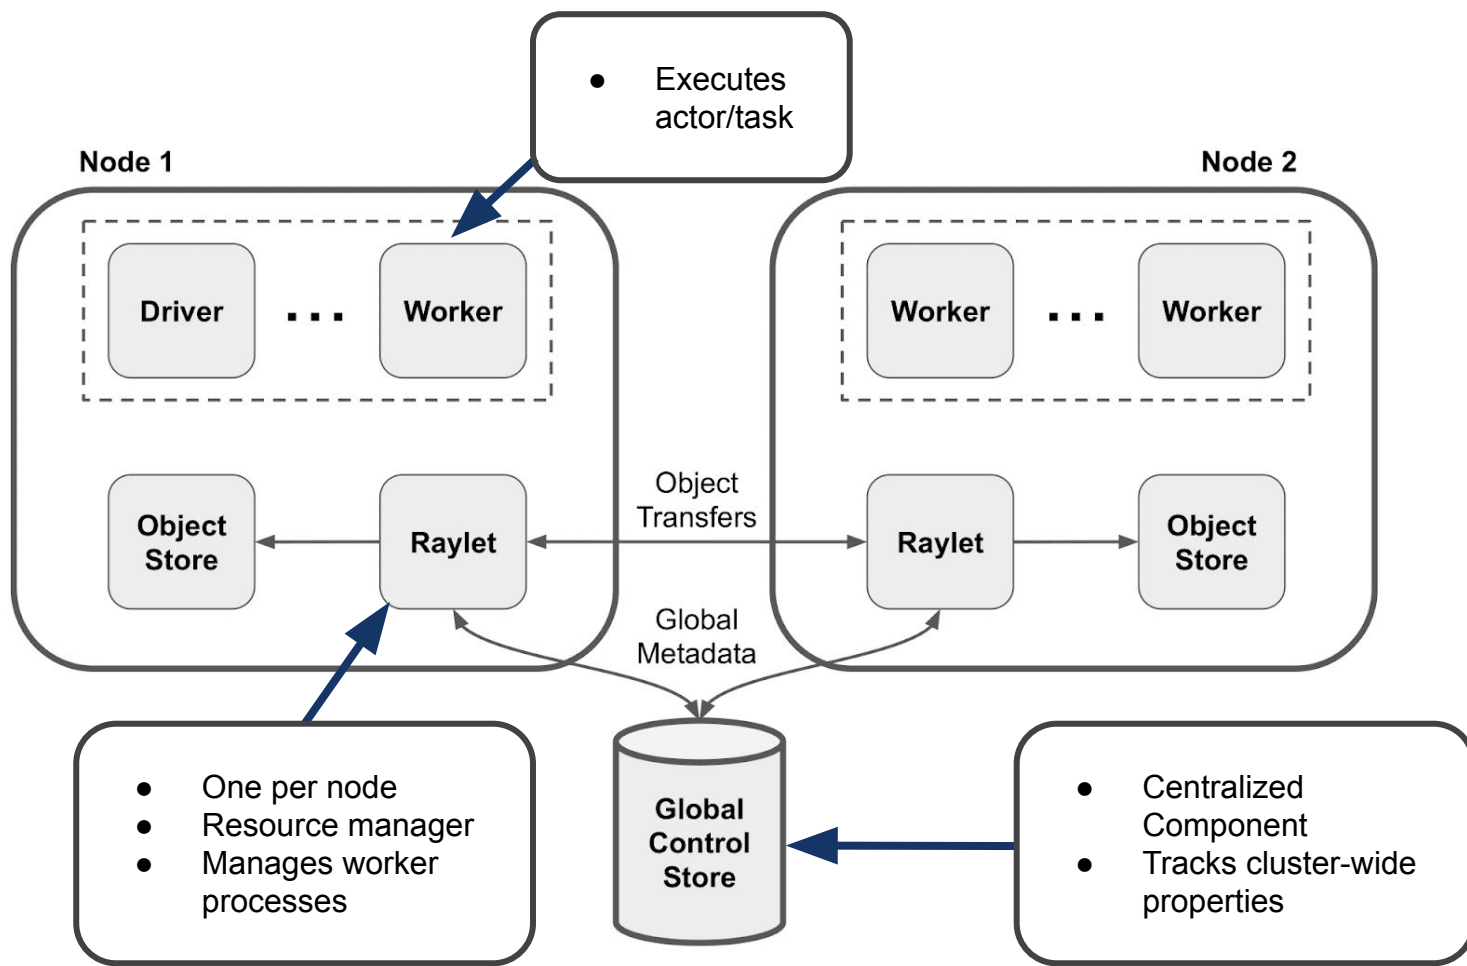


Ray Architecture & Components




An anatomy of a Ray cluster





Ray distributed design patterns & APIs



Ray Basic Design Patterns

→ Ray Parallel Tasks

- + Functions as stateless units of execution
- + Functions distributed across the cluster as tasks


→ Ray Objects as Futures

- + Distributed (immutable objects) store in the cluster
- + Fetched when materialized
- + Enable massive asynchronous parallelism

→ Ray Actors

- + Stateful service on a cluster
- + Enable Message passing

1. [Patterns for Parallel Programming](#)
2. [Ray Design Patterns](#)
3. [Ray Distributed Library Integration Patterns](#)



Python → Ray APIs




```
def f(x):  
    # do something with x:  
    y= ...  
    return y
```

Task




```
@ray.remote  
def f(x):  
    # do something with x:  
    y= ...  
    return y
```

Distributed



...



```
class Cls():  
    def __init__(self,  
x):  
    def f(self, a):  
        ...  
    def g(self, a):  
        ...
```

Actor




```
@ray.remote  
class Cls():  
    def  
__init__(self, x):  
    def f(self, a):  
        ...  
    def g(self, a):  
        ...
```

Distributed



...




```
import numpy as np  
a= np.arange(1, 10e6)  
b = a * 2
```

Distributed
immutable
object




```
import numpy as np  
a = np.arange(1, 10e6)  
obj_a = ray.put(a)  
b = ray.get(obj_a) * 2
```

Distributed



...



Function → Task

```
@ray.remote
def read_array(file):
    # read ndarray "a"
    # from "file"
    return a
```

```
@ray.remote
def add(a, b):
    return np.add(a, b)
```

```
id1 = read_array.remote(file1)
id2 = read_array.remote(file2)
id = add.remote(id1, id2)
sum = ray.get(id)
```

Class → Actor

```
@ray.remote(num_gpus=1)
class Counter(object):
    def __init__(self):
        self.value = 0
    def inc(self):
        self.value += 1
        return self.value
```


```
c = Counter.remote()
id4 = c.inc.remote()
id5 = c.inc.remote()
```

Task API

```
@ray.remote
def read_array(file):
    # read ndarray "a"
    # from "file"
    return a

@ray.remote
def add(a, b):
    return np.add(a, b)

id1 = read_array.remote(file1)
id2 = read_array.remote(file2)
id = add.remote(id1, id2)
sum = ray.get(id)
```

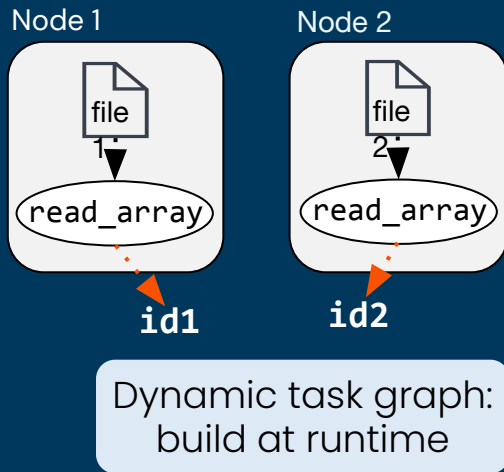


Task API

```
@ray.remote
def read_array(file):
    # read ndarray "a"
    # from "file"
    return a

@ray.remote
def add(a, b):
    return np.add(a, b)

id1 = read_array.remote(file1)
id2 = read_array.remote(file2)
id = add.remote(id1, id2)
sum = ray.get(id)
```




Task API

```
@ray.remote
def read_array(file):
    # read ndarray "a"
    # from "file"
    return a
```


```
@ray.remote
def add(a, b):
    return np.add(a, b)
```

```
id1 = read_array.remote(file1)
id2 = read_array.remote(file2)
id = add.remote(id1, id2)
sum = ray.get(id)
```

`ray.get()` block until
result available



Distributed Immutable object store



2
6

Spill over to external
storage

Distributed object store


Node 1

```
@ray.remote  
def f():  
    ...  
    return X
```

```
@ray.remote  
def g(a):  
    ...  
    return Y
```

```
id_X = f.remote()  
id_Y = g.remote(id_X)
```

Node 2




Shared object store

f()

id_X


Only X's id (`id_X`) is returned, not X's value

Distributed object store



`g(id_X)` is scheduled on same node, so X is never transferred

Examples of Distributed Applications with Ray



Distributed Applications with Ray

ML Libraries

- Ray AI Runtime
 - Ray Data, Tune, Train, Serve, RLlib
- Distributed scikit-learn/Joblib
- Distributed XGBoost on Ray
- Dask on Ray
- Modin on Ray

All using Ray design patterns

Monitoring Services

- WhyLabs
- Arize AI
- W & B

All using Ray design patterns

ML Platforms & Integrations

- Merlin (Shopify)
- Zero Copy (IBM)
- TorchX
- MLflow, Comet
- AirFlow
- HuggingFace
- Pycaret
- Ludwig AI


All using Ray design patterns

Ray Ecosystem: <https://docs.ray.io/en/latest/ray-overview/ray-libraries.html>

Key Takeaways

- Distributed computing is a necessity & norm
- Ray's vision: make distributed computing simple
 - + Don't have to be distributed programming expert
- Build your own disruptive apps & libraries with Ray
- Scale your ML workloads with Ray libraries (Ray AIR)

Let's go with 



Your Anyscale Cluster

- Console: <http://console.anyscale.com/>
- User name: <username@gmail.com>
- Password : password

console.anyscale.com/register?redirectTo=2

anyscale

Scale your application from your laptop to the cloud

Get started

Work email

Next

1

console.anyscale.com/register/login-user?org_id=org_1HxYtj@tEdauc580HyqthQET

anyscale

Scale your application from your laptop to the cloud

Log in

Email

yinhaonan55+200@gmail.com

Password

Forgot your password?

Log in

2

Your Anyscale Cluster

← → ↻ console.anyscale.com/o/ray-summit-tutorial-2022-test/clusters

anyscale

- Home
- Projects
- Interactive sessions
- Jobs
- Services
- Clusters**
- Configurations

Tutorial Account

Help

Feedback

< Collapse

Clusters

+ Create ▶ Start ⏸ Terminate 🗑 Delete

Search names Cluster status Created by Include archived

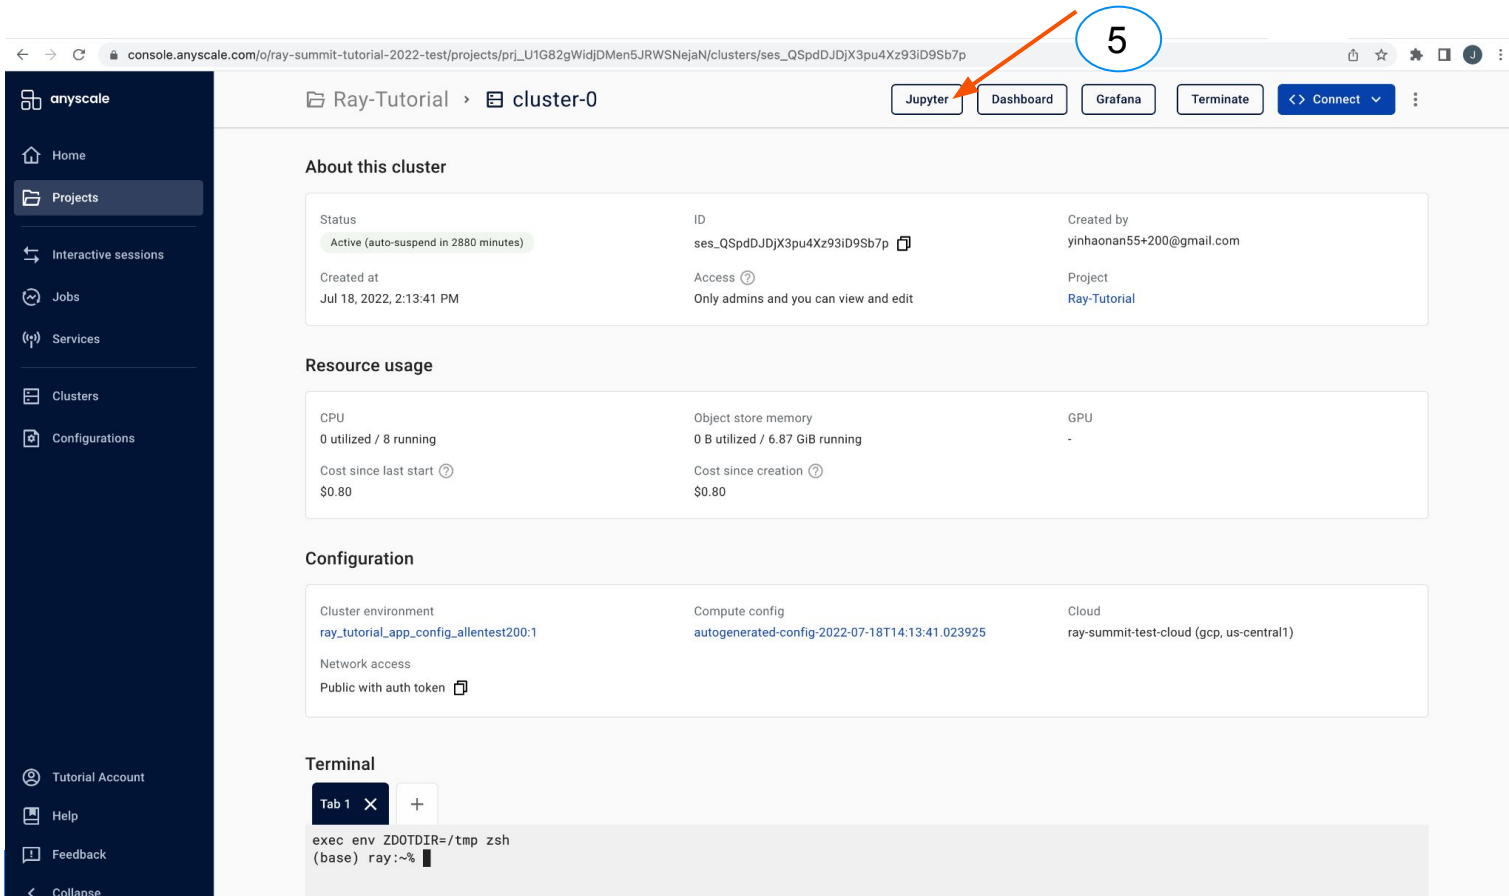
<input type="checkbox"/>	Name	Status ↓	Active resources	Cost ⓘ ↕	Cluster environment	Project	Cloud	Created by	Created
<input type="checkbox"/>	cluster-0	Terminated	None	\$0.80	ray_tutorial_app_config_allentest200:1	Ray-Tutorial	ray-summit-test-cloud (GCP)	Me	7/18

1 - 1 of 1 < >

3

4

Your Anyscale Cluster



The screenshot shows the Anyscale console interface. The browser address bar displays the URL: `console.anyscale.com/o/ray-summit-tutorial-2022-test/projects/prj_U1G82gWidjDMen5JRWsNejaN/clusters/ses_QSpdJDjX3pu4Xz93iD9Sb7p`. The page title is "Ray-Tutorial > cluster-0". The navigation sidebar on the left includes links for Home, Projects, Interactive sessions, Jobs, Services, Clusters, Configurations, Tutorial Account, Help, and Feedback. The main content area is divided into several sections:

- About this cluster**: A table showing cluster details.

Status	ID	Created by
Active (auto-suspend in 2880 minutes)	ses_QSpdJDjX3pu4Xz93iD9Sb7p	yinhaonan55+200@gmail.com
Created at	Access	Project
Jul 18, 2022, 2:13:41 PM	Only admins and you can view and edit	Ray-Tutorial
- Resource usage**: A table showing resource utilization.

CPU	Object store memory	GPU
0 utilized / 8 running	0 B utilized / 6.87 GiB running	-
Cost since last start	Cost since creation	
\$0.80	\$0.80	
- Configuration**: A table showing cluster configuration.

Cluster environment	Compute config	Cloud
ray_tutorial_app_config_allentest200:1	autogenerated-config-2022-07-18T14:13:41.023925	ray-summit-test-cloud (gcp, us-central1)
Network access		
Public with auth token		
- Terminal**: A terminal window with the following commands and output:

```
exec env ZDOTDIR=/tmp zsh
(base) ray:~% █
```

Your Anyscale Cluster

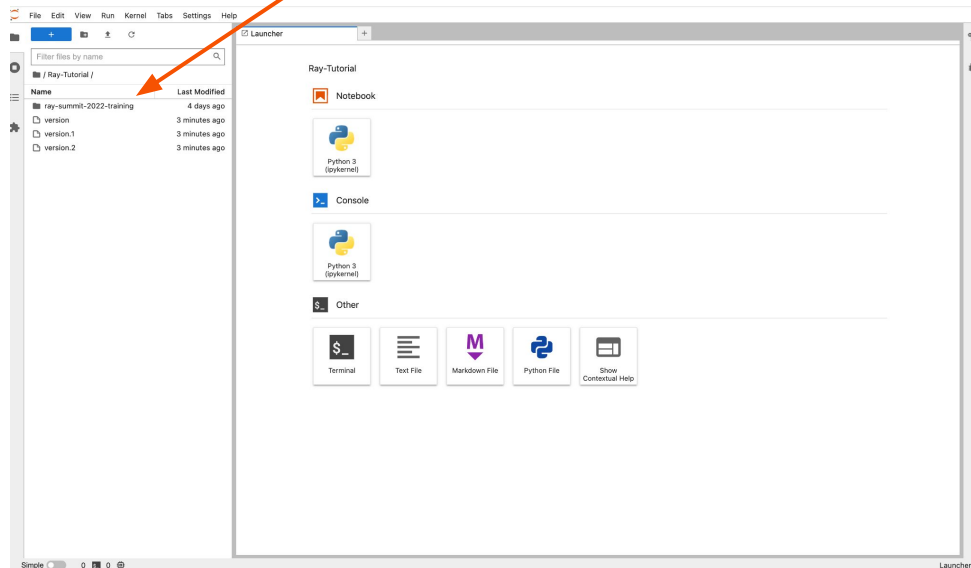
The screenshot displays the Anyscale JupyterLab interface. The browser address bar shows the URL: `dashboard-ses-qspddjdjx3pu4xz93id9sb7p.anyscale-hxyrj8r-g28hqzi-0000.anyscale-test-production.com/jupyter/lab`. The interface includes a top menu bar with options: File, Edit, View, Run, Kernel, Tabs, Settings, and Help. Below the menu is a toolbar with icons for file operations. On the left, a sidebar shows a file explorer with a search bar labeled "Filter files by name". A table lists files with columns "Name" and "Last Modified":

Name	Last Modified
anaconda3	2 months ago
Ray-Tutorial	2 minutes ago
requireme...	2 months ago

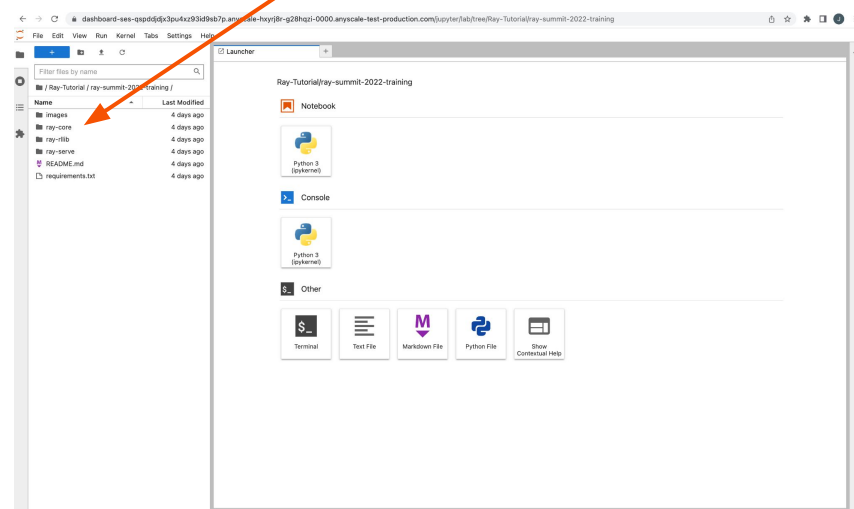
An orange arrow points to the "anaconda3" entry. The main area is titled "Launcher" and features a large blue circle with the number "6" in the top-left corner. Below this, the "Notebook" section contains a "Python 3 (pykernel)" icon. The "Console" section also contains a "Python 3 (pykernel)" icon. The "Other" section displays a row of icons: Terminal, Text File, Markdown File, Python File, and Show Contextual Help. The bottom status bar shows "Simple" and "Launcher" tabs.

Your Anyscale Cluster

7



8



Thank you.

Tell us what you think...

<https://bit.ly/ray-core-summit2022>




Tell us what you think...

<https://bit.ly/ray-core-summit2022>

Here is the Title Slide

Firstname Lastname, Company

Colors



Here is a basic Dark Slide

Here is a Basic Light Slide

How about a slide with 2 options?

Here is an info card

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation

Here is an info card

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation

How about a slide with 3?

Here is an info card

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation

Here is an info card

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation

Here is an info card

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation

Here is a Section Header



Here is a Section Header



Here is a Section Header



Thank you.

Follow up information can go here.

